

LSM_ TREES

A DEEP DIVE INTO THE DS POWERING
MODERN DATA PIPELINES

UC BERKELEY EXTENSION - DSA

A PRESENTATION BY REZA RAJAN

MOTIVATION

I have selected this topic to deepen my understanding into how modern databases are designed to handle the data of today.

Without LSM Trees, believe it or not, our day-to-day lives will look much different.

INTRODUCTION

THE PROBLEM

Traditional storage systems were not designed for such high-frequency workloads.

Random disk I/O was becoming a huge bottleneck.

THE PROBLEM

Engineers needed a way to handle this constant flood of incoming data efficiently, without slowing down reads or overwhelming the storage system.

In comes LSM Trees.

TOPICS OF DISCUSSION

IN THIS
DEEP DIVE...


Design principles

Implementation

Scope and Limitations

Big O Analysis

Demo



DESIGN PRINCIPLES

LSM TREES

DEFINITION:


LSM Tree = Log-Structured Merge Tree.

An abstract data structure designed to handle workloads with high-frequency writes by buffering updates in memory and periodically merging them into sorted, immutable files on disk.

DESIGN PRINCIPLES

- Leverage fast sequential I/O
- Decouple write layer and read layer
- Utilize high-speed memory, and flush data to disk over time for durability
- Keep data on disk sorted for fast lookups


LSM Trees



HOW IT WORKS - PROCESS OVERVIEW

We can broadly materialize these principles into three processes: collect, flush and compact.

LSM Trees



HOW IT WORKS - COLLECTION

In the LSM Tree, all operations are treated as append-only writes—including inserts, updates, and deletes. They are first collected in the **WAL**, and then pushed to a **Memtable**.

Key Concept: WAL and Memtable form the high-throughput “write” layer of the LSM Tree.

WRITE-AHEAD LOG (WAL)

The Write-Ahead Log is an **append-only** log.

Supports sequential I/O for fast operation.

Resides on disk to ensure durability in the event of a crash.


MEMTABLE

In-memory data store, which *may* be sorted.

Reads and writes benefit from **fast in-memory storage**.

Acts as caching layer for recent data.

LSM Trees



HOW IT WORKS - COLLECTION

In the LSM Tree, all operations are treated as append-only writes—including inserts, updates, and deletes. They are first collected in the **WAL**, and then pushed to a **Memtable**.

Key Concept: WAL and Memtable form the high-throughput “write” layer of the LSM Tree.

WRITE-AHEAD LOG (WAL)

The Write-Ahead Log is an **append-only** log.

Supports sequential I/O for fast operation.

Resides on disk to ensure durability in the event of a crash.

MEMTABLE

In-memory data store, which *may* be sorted.

Reads and writes benefit from **fast in-memory storage**.


Acts as caching layer for recent data.

CAVEAT

Memory is a limited resource.

What happens when the Memtable becomes full?

LSM Trees



HOW IT WORKS - FLUSHING


Once the Memtable is full, its data is persisted to disk as **SSTables**, and the **Memtable is freed to accept more writes**.

Key Concept: SSTables store sorted, structured data on disk. They support efficient reads and form the “read” layer of the LSM Tree.

SSTABLES

Immutable, sorted files on disk. Provide a structured, queryable format for efficient reads (e.g., Binary Search). Each flush creates a **new** SSTable, skipping in-place updates and leveraging sequential I/O.

LSM Trees



HOW IT WORKS - FLUSHING

Once the Memtable is full, its data is persisted to disk as **SSTables**, and the **Memtable is freed to accept more writes**.

Key Concept: SSTables store sorted, structured data on disk. They support efficient reads and form the “read” layer of the LSM Tree.

SSTABLES

Immutable, sorted files on disk. Provide a structured, queryable format for efficient reads (e.g., Binary Search). Each flush creates a **new** SSTable, skipping in-place updates and leveraging sequential I/O.


CAVEAT

Over time, the number of SSTable files accumulate from repeated flushes.

Scans must traverse all these files to find a key, slowing read operations with each flush of the Memtable.

This problem is called **read amplification**.

LSM Trees



HOW IT WORKS - COMPACTATION

Compaction counteracts read-amplification by reducing the sparsity of SSTable files on disk.

Involves organizing SSTables into logical **levels** through merging.

Merging maintains sorted order, prunes outdated and deleted entries, and reduces the total number of SSTable files at each level.

LEVEL 0

The first level.

Stores data **directly flushed** from the Memtable.


LEVEL 1

All SSTables from Level 0 are **merged into a single, larger, but more compact** SSTable in Level 1, which is the next level down.

Further Levels

This process happens at each level, reducing the sparsity of SSTables each time.

LSM Trees



HOW IT WORKS - SSTABLE FILTERS

Despite their sorted nature, there is no way to determine if a key exists within an SSTable—or if we have the latest version—without scanning all SSTables until it is found.


Two scenarios:

Range Scans may require scanning across multiple tables to find multiple values.

Scanning for non-existent data requires scanning all tables before confirming a key does not exist.

Other data structures, like B-Trees, do not suffer from this problem as their data is not stored across multiple files like with LSM Trees.

LSM Trees



HOW IT WORKS - SSTABLE FILTERS

Scanning for **non-existent** data requires scanning **all** tables.

Wouldn't it be nice if we could just tell if a key does not exist in a table? That would make the operation **constant time** instead of **linear**.


We can — by using Bloom Filters!

BLOOM FILTERS

Bloom filters
probabilistic data
structure.

Quickly indicate
whether a key **may**
exist in a given
SSTable or definitely
does not exist.

LSM Trees



HOW IT WORKS - SSTABLE FILTERS

Range Scans may require scanning across multiple tables.

Wouldn't it be nice if we could just seek to a value in an SSTable?


We can (sort of) — by using Sparse Indexes!

SPARSE INDEXES

Sparse indexes point to the **approximate location** of a key within an SSTable.

Enables **fast, targeted access** to keys, drastically reducing search time.

LSM Trees



HOW IT WORKS - SSTABLE FILTERS

Together, Bloom Filters and Sparse Indexes help us overcome the inefficiencies with reading across multiple SSTable files.

BLOOM FILTERS


Bloom filters probabilistic data structure.

Quickly indicate whether a key **may exist** in a given SSTable or **definitely does not exist**.

SPARSE INDEXES

Sparse indexes point to the **approximate location** of a key within an SSTable.

Enables **fast, targeted access** to keys, drastically reducing search time.



SCOPE AND LIMITATIONS



SCOPE AND LIMITATIONS

SCOPE


LSM Trees are:

- **Designed for sustained, high-volume writes with durable storage;** not for workloads dominated by point reads or in-memory caching.
- **Built to handle datasets far larger than available RAM;** not for data that fits entirely in memory.

Workloads must justify their requirements for high write throughput and data volume before even considering LSM Trees.

LIMITATIONS

- There is inherent increased disk I/O when performing compactions; schedule compactions wisely.
- Without proper SSTable management policies, read latency grow unbounded as the number of tables increase.
- Consider extra space needed for compaction when selecting storage.



CODE IMPLEMENTATION



IMPLEMENTATION - BACKENDS

WAL

WAL implements the append method.

append accepts a key and value, as well as timestamp, known as a record. It writes the record to disk in a binary format.

It encodes the type of operation in an op code; either PUT or DELETE.

In the event of a crash, this log is replayed to recover the state of the Memtable.

```
def append(self, key: Key, value: Value | None, ts: Timestamp) -> int:  
    """Append a record to WAL.  
  
    Args:  
        key: Binary key  
        value: Binary value or None for tombstone  
        ts: Monotonic timestamp  
  
    Returns:  
        WAL sequence number  
    """  
    if self._fd is None:  
        raise RuntimeError("WAL is closed") # noqa: TRY003  
  
    op_code = OP_DELETE if value is None else OP_PUT  
    value_bytes = value if value is not None else b"""  
  
    # Build record  
    key_len = len(key)  
    value_len = len(value_bytes)  
  
    # Pack payload (without CRC)  
    payload = struct.pack("<I", MAGIC)  
    payload += struct.pack("<Q", key_len)  
    payload += key  
    payload += struct.pack("<Q", value_len)  
    payload += value_bytes  
    payload += struct.pack("<Q", ts)  
    payload += struct.pack("<B", op_code)  
  
    # Calculate CRC32 of payload  
    crc = zlib.crc32(payload)  
    record = payload + struct.pack("<I", crc)  
  
    # Write to file  
    self._fd.write(record)  
    if self.flush_every_write:  
        self._fd.flush()  
        os.fsync(self._fd.fileno())  
  
    self.sequence += 1  
    logger.debug(f"Appended record seq={self.sequence}, key_len={key_len}, ts={ts}")  
  
    return self.sequence
```



IMPLEMENTATION - BACKENDS

MEMTABLE

Memtable implements the **put**, **delete**, **get** and **clear** methods.

put accepts a record and appends it to the Memtable in **sorted order**.

delete accepts a key to delete and sets its value to None.

get returns a key and its associated value if the key exists; otherwise it returns None.

clear deletes all Memtable entries, freeing it for further writes.

Design Decision

Either sort entries on write, or sort them on flush.

This implementation uses `SortedDict` to sort entries on write.

```
class SimpleMemtable:  
    """In-memory sorted structure holding recent writes.  
  
    Maintains key-value pairs with timestamps in sorted order.  
    Supports efficient range queries and size tracking.  
  
    Invariants:  
        - Keys are always maintained in sorted order  
        - Most recent value per key (by timestamp) is accessible  
        - Size includes approximate overhead of data structures  
    """  
  
    def __init__(self):  
        """Initialize empty memtable.  
        self._data: SortedDict = SortedDict()  
        self._size_bytes: int = 0  
  
    def put(self, key: Key, value: Value, ts: Timestamp) -> None:  
        """Insert or update key with value and timestamp.  
        # Track size delta  
        if key in self._data:  
            old_value, _old_ts = self._data[key]  
            old_size = len(key) + (len(old_value) if old_value else 0) + 8  
            self._size_bytes -= old_size  
  
            self._data[key] = (value, ts)  
            new_size = len(key) + len(value) + 8 # key + value + timestamp  
            self._size_bytes += new_size  
  
        def delete(self, key: Key, ts: Timestamp) -> None:  
            """Mark key as tombstone with timestamp.  
            # Track size delta  
            if key in self._data:  
                old_value, _old_ts = self._data[key]  
                old_size = len(key) + (len(old_value) if old_value else 0) + 8  
                self._size_bytes -= old_size  
  
                self._data[key] = (None, ts)  
                new_size = len(key) + 8 # key + timestamp (no value)  
                self._size_bytes += new_size  
  
        def get(self, key: Key) -> tuple[Value | None, Timestamp] | None:  
            """Return (value_or_none, timestamp) if key found; else None.  
            # If value_or_none is None, it is a tombstone.  
            #  
            return self._data.get(key)
```



IMPLEMENTATION - BACKENDS

SSTABLE

SSTable implements the add, may_contain, and get methods.

add is only used when flushing Memtables to disk.

It reads a record from the Memtable, serializes it into a binary format, and then appends it to the SSTable on disk.

add also ensures sorted order.

Since we push the sorting process to the Memtable, this operation runs in **constant time**.

Finally, each add operation creates a Bloom Filter entry, and potentially one for the Sparse Index.

```
def add(self, key: Key, value: Value | None, ts: Timestamp) -> None:  
    """Append record to the writer (must be added in sorted order)."""  
    # Verify sorted order  
    if self._last_key is not None and key <= self._last_key:  
        raise SSTableError("Unsorted keys") # noqa: TRY003  
  
    # Record current offset  
    if self._fd is None:  
        msg = "Writer closed"  
        raise SSTableError(msg)  
    fd = self._fd  
    offset = fd.tell()  
  
    # Update metadata  
    if self._min_key is None:  
        self._min_key = key  
    self._max_key = key  
  
    if self._min_ts is None or self._max_ts is None:  
        self._min_ts = ts  
        self._max_ts = ts  
    else:  
        self._min_ts = min(self._min_ts, ts)  
        self._max_ts = max(self._max_ts, ts)  
  
    # Sample for sparse index  
    if self._count % self.index.interval == 0:  
        self._index.append((key, offset))  
  
    # Collect for bloom filter  
    self._keys_for_bloom.append(key)  
  
    # Write record  
    value_bytes = value if value is not None else b""  
    key_len = len(key)  
    value_len = len(value_bytes)  
  
    fd.write(struct.pack("<Q", key_len))  
    fd.write(key)  
    fd.write(struct.pack("<Q", value_len))  
    fd.write(value_bytes)  
    fd.write(struct.pack("<Q", ts))  
  
    self._count += 1  
    self._last_key = key
```



IMPLEMENTATION - BACKENDS

SSTABLE

SSTable implements the add, may_contain, and get methods.

may_contain checks the Bloom Filter for a key.

It returns True if a key may exist, and False otherwise.

```
def may_contain(self, key: Key) -> bool:  
    """Use Bloom filter to test potential presence."""  
    # Check key range first  
    if self._min_key and key < self._min_key:  
        return False  
    if self._max_key and key > self._max_key:  
        return False  
    return key in self._bloom
```

IMPLEMENTATION - BACKENDS



SSTABLE

SSTable implements the add, may_contain, and get methods.

get returns the value of an associated key if the Bloom Filter passes, and utilizes the Sparse Index to quicken lookups.

Binary Search

Read operations capitalize on the sorted nature of SSTables, utilizing **binary search** for **logarithmic** lookups.

```
def get(self, key: Key) -> tuple[Value | None, Timestamp] | None:  
    """Return (value_or_none, ts) or None if key definitely not present."""  
    if not self.may_contain(key):  
        return None  
  
    self._ensure_open()  
  
    # Binary search in sparse index  
    offset = self._find_block_offset(key)  
  
    # Scan from offset  
    if self._fd is None:  
        self._ensure_open()  
        assert self._fd is not None  
        self._fd.seek(offset)  
  
    while True:  
        # Try to read a record  
        key_len_bytes = self._fd.read(8)  
        if len(key_len_bytes) < 8:  
            break # EOF  
  
        record_key_len = struct.unpack("<Q", key_len_bytes)[0]  
        record_key = self._fd.read(record_key_len)  
  
        if len(record_key) < record_key_len:  
            break # Corrupted  
  
        # If we've passed the key, it doesn't exist  
        if record_key > key:  
            break  
  
        # Read value and timestamp  
        value_len = struct.unpack("<Q", self._fd.read(8))[0]  
        value_bytes = self._fd.read(value_len)  
        ts = struct.unpack("<Q", self._fd.read(8))[0]  
  
        if record_key == key:  
            value = value_bytes if value_len > 0 else None  
            return (value, ts)  
  
    return None
```

IMPLEMENTATION - BACKENDS



COMPACTOR

The compactor takes in a sequence of SSTables at a given level and merges them into a single SSTable.

While merging, it maintains sorted order, and prunes deleted or outdated records.

```
def compact(
    self, input_tables: Sequence[SSTableMeta], target_level: int
) -> Sequence[SSTableMeta]:
    """Perform merge compaction and write output SSTables.

    Args:
        input_tables: SSTables to compact
        target_level: Target level for output

    Returns:
        List of produced SSTable metadata
    """
    if not input_tables:
        return []

    logger.info(f"Compacting {len(input_tables)} SSTables to level {target_level}")

    # Open all input SSTables
    readers: List[SSTableReader] = []
    for meta in input_tables:
        reader: SSTableReader = SimpleSSTableReader(meta["data_path"], meta["meta_path"])
        readers.append(reader)

    try:
        # Merge all iterators
        merged = self._merge_iterators(readers)

        # Write to output SSTables
        output_metas = self._write_output(merged, target_level)

        logger.info(f"Compaction produced {len(output_metas)} SSTables")
        return output_metas

    finally:
        # Close all readers
        for reader in readers:
            reader.close()
```

IMPLEMENTATION - BACKENDS

COMPACTOR

This implementation uses **min heaps** under-the-hood to maintain sorted order across SSTables.

Note that due to implementation specifics, the overall time complexity for compaction is **approximately linear - not quasilinear** as you may expect for heaps.

```
def _merge_iterators(self, readers: Sequence[SSTableReader]) -> Iterator[Record]:
    """Merge multiple sorted iterators, keeping newest by timestamp."""
    # Build heap of (key, -ts, value, reader_idx, iterator)
    heap = []
    iterators = [reader.iter_range(None, None) for reader in readers]

    # Initialize heap
    for idx, it in enumerate(iterators):
        try:
            key, value, ts = next(it)
            heapq.heappush(heap, (key, -ts, value, idx, it))
        except StopIteration: # noqa: PERF203
            pass

    last_key = None
    last_value = None
    last_ts = None

    while heap:
        key, neg_ts, value, idx, it = heapq.heappop(heap)
        ts = -neg_ts

        # Advance iterator
        try:
            next_key, next_value, next_ts = next(it)
            heapq.heappush(heap, (next_key, -next_ts, next_value, idx, it))
        except StopIteration:
            pass

        # Emit if key changed and we have a value to emit
        if last_key is not None and key != last_key:
            # Check if tombstone should be kept
            assert last_ts is not None
            if last_value is not None or self._should_keep_tombstone(last_ts):
                yield (last_key, last_value, last_ts)
            last_key = None

        # Update last seen (highest timestamp wins)
        if last_key is None or key != last_key:
            last_key = key
            last_value = value
            last_ts = ts
        elif ts > last_ts:
            last_value = value
            last_ts = ts

    # Emit final record
    if (
        last_key is not None
        and (
            last_value is not None
            or (last_ts is not None and self._should_keep_tombstone(last_ts))
        )
    ):
        assert last_ts is not None
        yield (last_key, last_value, last_ts)
```

BIG_O ANALYSIS

BIG_O ANALYSIS

The WAL has two operations: **append** and **replay**.

Append: Add entry to log

Replay: Reads the log from start to end.

	Best Case	Average Case	Worst Case
Append	$O(1)$	$O(1)$	$O(1)$
Replay	$O(n)$	$O(n)$	$O(n)$
WRITE-AHEAD LOG (WAL)			
Append operations run in constant time .			
Replay operations run in linear time .			

BIG_O ANALYSIS

The Memtable has five operations: **put**,
delete, **get**, **scan** and **flush**.

	Best Case	Average Case	Worst Case
MEMTABLE Implementation uses SortedDict - average logarithmic time for all operations. Flush must pop all elements - linear time.	Put $O(1)$	$O(\log n)$	$O(\log n)$
	Delete $O(1)$	$O(\log n)$	$O(n)$
	Get $O(1)$	$O(\log n)$	$O(n)$
n - total number of records k - range of records to scan	Scan —	$O(k \log n)$	$O(n * k)$
	Flush —	$O(n)$	$O(n)$

BIG_O ANALYSIS

The SSTable has three operations: **point read**, **range read** and **compact**.

SSTables

Reads employ **Binary Search**, therefore **logarithmic time**.

Search against Sparse Indexes + Bloom Filter


n - number of SSTables

k - number of records in an SSTable

b - SSTable block size

I - number of index entries in the sparse index (< k)

	Best Case	Average Case	Worst Case
Point	$O(1)$	$O(\log I + b)$	$O(\log I + b)$
Range	$O(1)$	$O(\log I + b)$	$O(\log I + b)$
Compact	—	$O(n*k)$	$O(n*k)$



DEMO

DEMO - CLI




```
>>> Pu█
```

Interacting with the LSM Tree via Python CLI




DEMO - BENCHMARKS

SYNCHRONOUS VS ASYNCHRONOUS COMPACTION



Synchronous Compaction



Asynchronous Compaction

Memtable size reduced to force frequent flushes and compactions




DEMO - BENCHMARKS

SYNCHRONOUS VS ASYNCHRONOUS COMPACTION



Synchronous Compaction




Asynchronous Compaction

Memtable size reduced to force frequent flushes and compactions


DEMO - BENCHMARKS



SYNCHRONOUS VS ASYNCHRONOUS COMPACTION




Synchronous Compaction



Asynchronous Compaction

Memtable size reduced to force frequent flushes and compactions

IMPROVEMENTS

The background of the slide features a subtle, abstract design composed of overlapping squares. There are several large, semi-transparent squares in a light teal or mint green color, which are partially obscured by smaller, solid white squares. These shapes are scattered across the slide, creating a modern and clean aesthetic.

1. Block Cache

Introduce an **LRU-based cache** for frequently accessed blocks or index pages. This could drastically reduce SSTable read latency.

2. Concurrency

Apply **MVCC** to allow non-blocking reads during ongoing writes, flushes, or compactions.

3. Compaction Policies

Add adaptive policies (size-tiered, hybrid, or leveled with configurable thresholds) to fit different performance targets.

THANK YOU!